# black hat EUROPE 2024

**DECEMBER 11-12, 2024** BRIEFINGS

# Improving Side-Channel Protections for Intel® TDX

\*Scott Constable | Defensive Security Researcher, Intel Labs, Intel Corporation \*Nagaraju (Raju) Kodalapura Principal Engineer, Offensive Security Research, Intel Corporation Baruch Chaikin Principal Engineer, CPU Architecture, Intel Corporation



#### Acknowledgements

- Intel Labs
- Intel Product Security and Assurance (IPAS) research
- Intel CPU Security Effectiveness Team
- Intel CPU Architecture Team
- The TDXdown researchers at TU Lübeck: Luca Wilke, Florian Sieck, Thomas Eisenbarth



Scott Constable



Nagaraju (Raju) Kodalapura



**Baruch Chaikin** 



#### **Agenda**

- Intro to Confidential Computing, Intel TDX (Trust Domain Extensions), side-channel Attacks, and malicious single-stepping
- Pre-TDX PoC (Proof of Concept) TDX-step exploit and mitigation
- Techniques to bypass the TDX-Step mitigation, and intro to the new ICSSD (Instruction Counting Single-Step Defense) feature
- Comparison with the SGX-Step mitigation



#### What is Confidential Computing (CC)?



Protects data at rest (in storage, a database, etc.).

- Data encryption
- Access control



Protects data in transit (over a network, PCI bus, etc.)

- HTTPS
- TLS

The Focus of CC

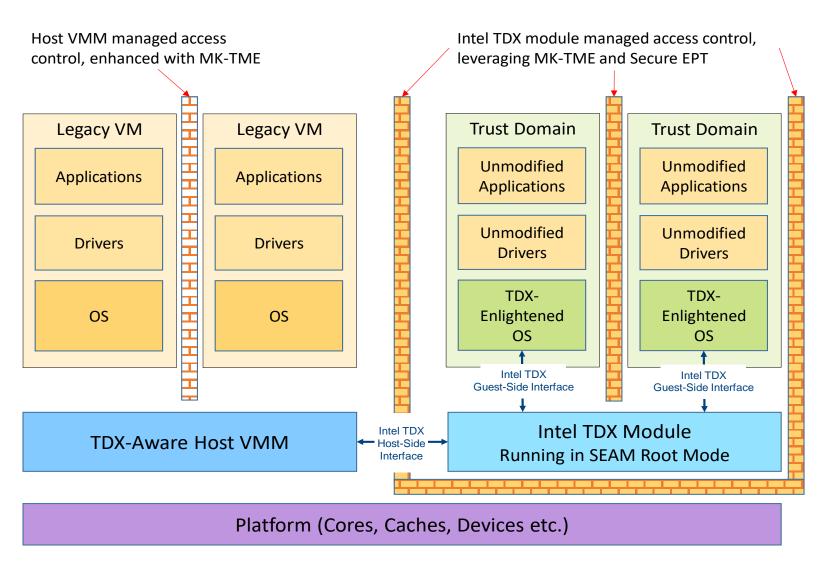


Protects data in use (within a CPU, XPU, etc.)

 Hardware-based, attested Trusted Execution Environments (TEEs) such as Intel TDX and Intel SGX



#### What is Intel® TDX?



Intel TDX is a CC technology that provides confidentiality and integrity for data in use by tenant VMs, called <u>Trust Domains</u> (TDs)

- Objective: Remove the Virtual Machine Monitor (VMM) and other system SW from the TDs' TCB
- TDX Module: Intel-signed security services module responsible for enforcing security policies for TDs
- SEAM: A new Secure Arbitration Mode (SEAM) hosts the TDX module



#### **TDX Threat Model**

**TRUSTED BY TD** 

INTEL® TDX MODULE

INTEL AUTHENTHICATED CODE MODULES (ACM)

TD QUOTING ENCLAVE

**INTEL CPU HARDWARE** 

Some side-channel attacks are out of scope for the TDX threat model

**NOT TRUSTED BY TD** 

**PLATFORM ADMIN** 

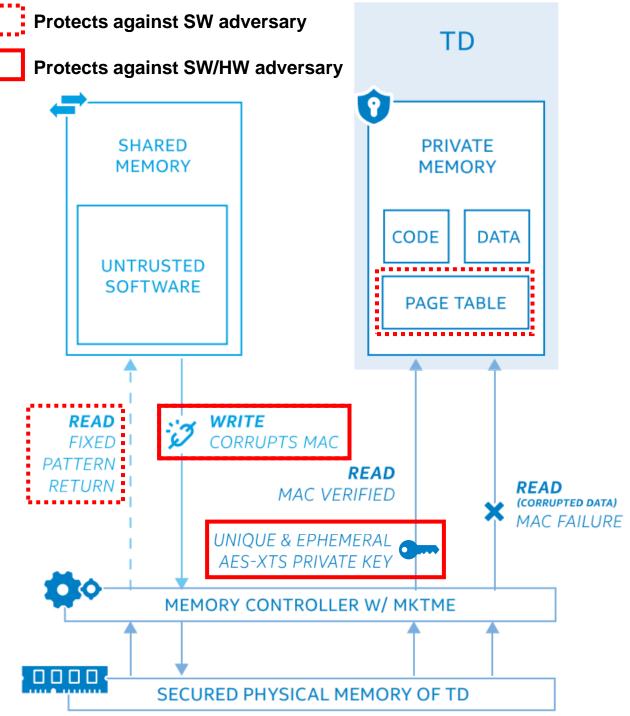
DISCRETE AND INTEGRATED DEVICES

**ALL OTHER SOFTWARE** 

OTHER PLATFORM FIRMWARE

**HOST-OS/VMM** 

BIOS/SMM





#### What is a Side Channel?

Victim Context

Do a cat thing! Do something else! Do a different cat thing!



#### **Attacker Context**





Do a teapot thing! My teapot thing was *slow*, so the victim must have done a cat thing! Do a teapot thing! My teapot thing was fast, so the victim must *not* have done a cat thing! Do a teapot thing! My teapot thing was *slow*, so the victim must have done a cat thing!





### Adversary's Challenge: Shutter Speed



just triggers a single

operation at a time,

side-channel

measurement?

What if the adversary and then performs a TD runs

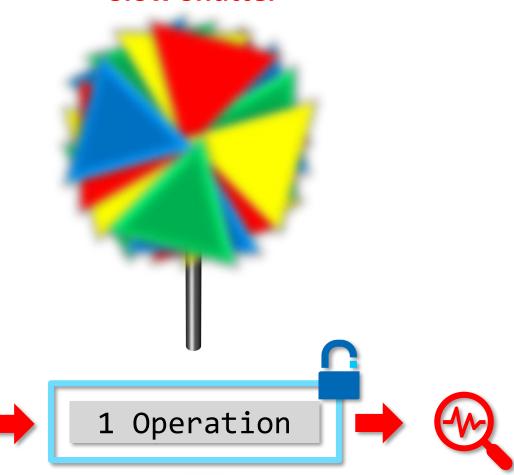
TD instruction stream (typically 1.8 GHz or more):



## Adversary's Challenge: Shutter Speed 🔯



**Slow Shutter** 

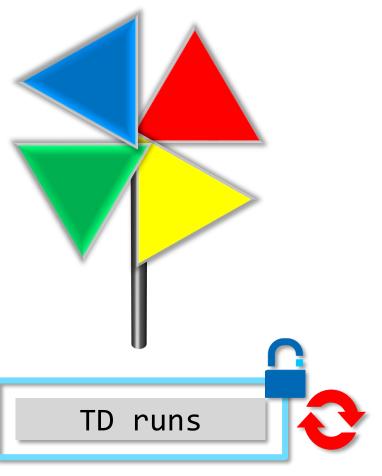


TD instruction stream (typically 1.8 GHz or more):



### Adversary's Challenge: Shutter Speed 🔯





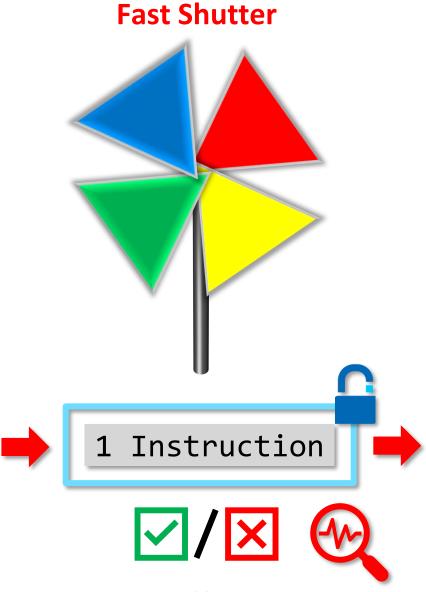
Let's try something more precise: can the adversary trigger a single *instruction* at a time?

TD instruction stream (typically 1.8 GHz or more):



### Adversary's Challenge: Shutter Speed





This technique was pioneered in a framework called:





#### **Searching for an Ideal HW Primitive**

#### Intel® 64 and IA-32 Architectures Software Developer's Manual

## CHAPTER 12 ADVANCED PROGRAMMABLE INTERRUPT CONTROLLER (APIC)

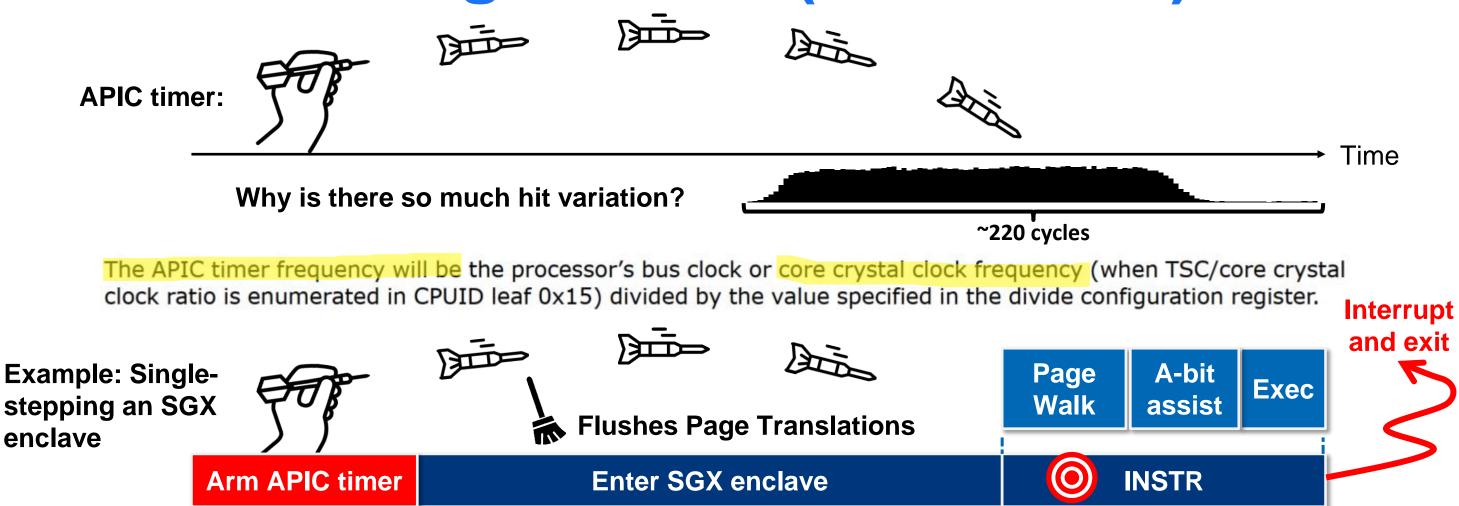
#### 12.5.4 APIC Timer

The local APIC unit contains a 32-bit programmable timer that is available to software to time events or operations. This timer is set up by programming four registers: the divide configuration register (see Figure 12-10), the initial-count and current-count registers (see Figure 12-11), and the LVT timer register (see Figure 12-8).

TSC-deadline mode allows software to use the local APIC timer to signal an interrupt at an absolute time. In TSC-deadline mode, writes to the initial-count register are ignored; and current-count register always reads 0. Instead, timer behavior is controlled using the IA32\_TSC\_DEADLINE MSR.



#### A Small Target for an (Inaccurate) APIC



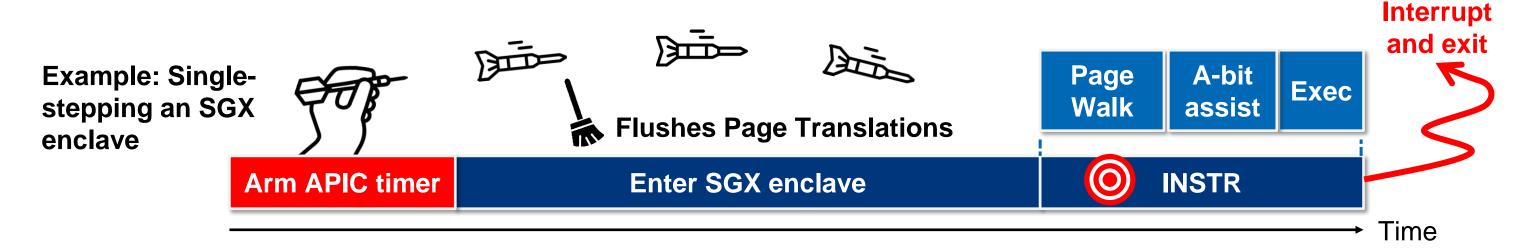
Time



#### A Small Target for an (Inaccurate) APIC

Key Idea 1: Don't try to hit a small target! Instead, make the target *slower* and therefore

## bigger





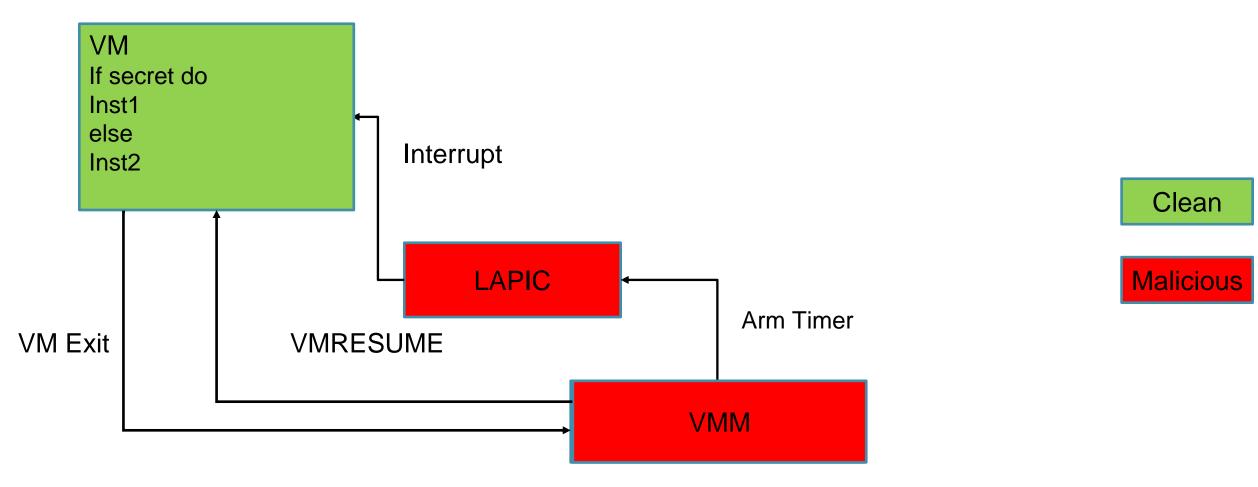
#### **Agenda**

- Intro to Confidential Computing, Intel TDX (Trust Domain Extensions), side-channel Attacks, and malicious single-stepping
- Pre-TDX PoC (Proof of Concept) TDX-step exploit and mitigation
- Techniques to bypass the TDX-Step mitigation, and intro to the new ICSSD (Instruction Counting Single-Step Defense) feature
- Comparison with the SGX-Step mitigation



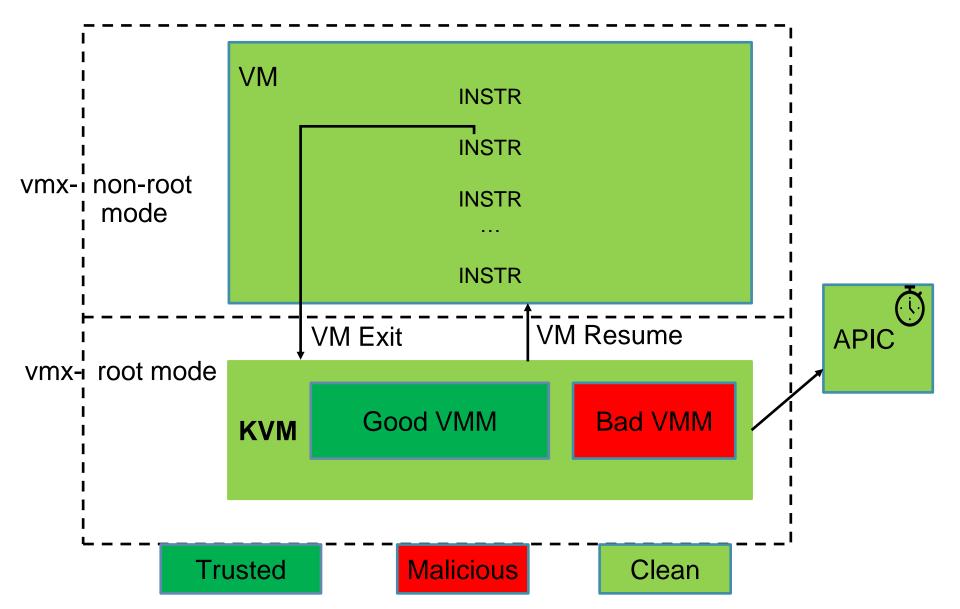
#### **TDX-Step Attack Scenario**

- PoC (Proof of Concept) Exploit was developed in a virtualization setup (no TDX support)
- Interrupt based attack: VMM mounts attack on VM/TD application using APIC





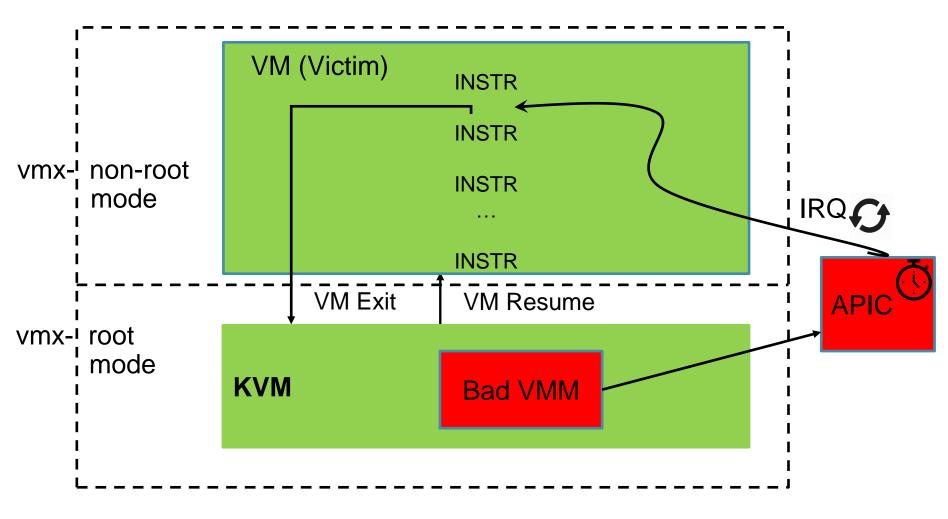
#### **Pre-TDX Test Setup**



- KVM (Host VMM) as the host used as the both Good (Mitigation) and Bad (Attack) Actor
- Ubuntu VM (Virtual Machine) as the Guest, running "strlen" workload/app as a victim app
- Good VMM portion of the KVM code was eventually implemented in Intel TDX Module, which is responsible for mitigating the TDX-Step attack



#### **TDX-Step Attack Realized**



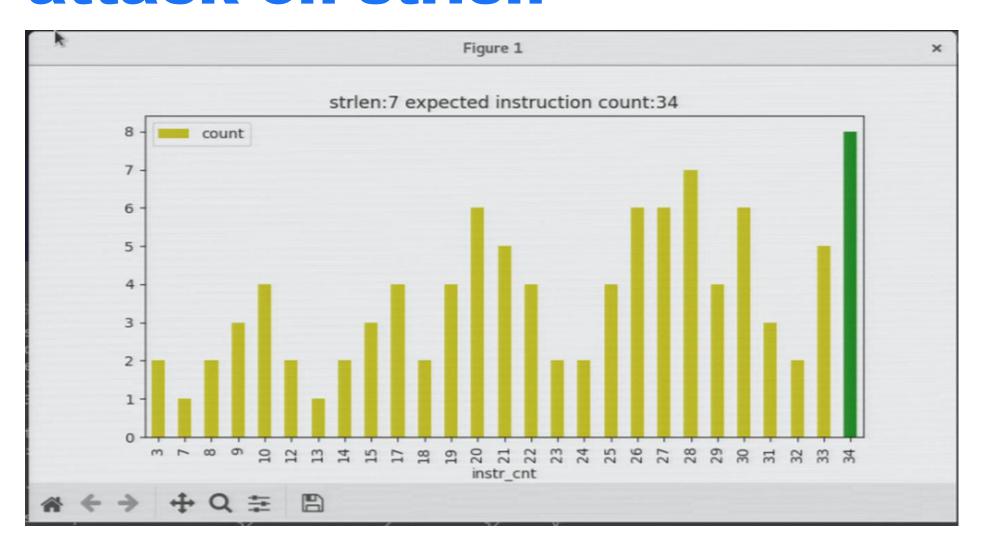
Clean

Malicious

- Bad VMM configures the APIC timer to cause periodic interrupts to Victim VM on an instruction boundary
- Once interrupted, victim VM will exit to the host through VM Exit
- Now Bad VMM has the ability conduct side channel analysis on the VM instruction execution
- Also, bad VMM will cause VM
   Resume to make the victim VM
   to execute all the instructions



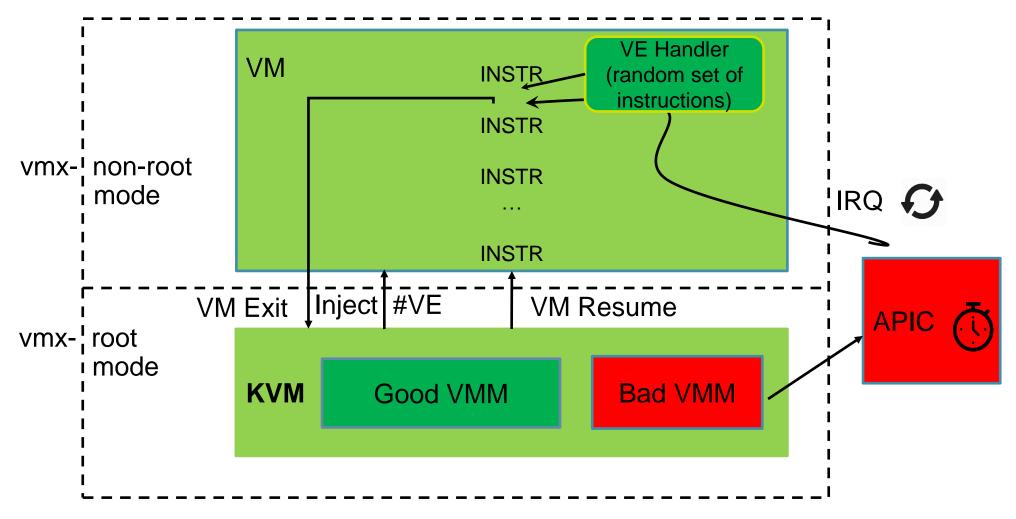
# Proof of Concept attack: Interrupt-based attack on strien



- Victim VM executing the strlen 7 consisted of 34 instructions
- Out of 100 experiments, the PoC attack was able to cause single-step attack successfully for 8 times
- This was sufficient to prove that TDX-Step attack indeed is real

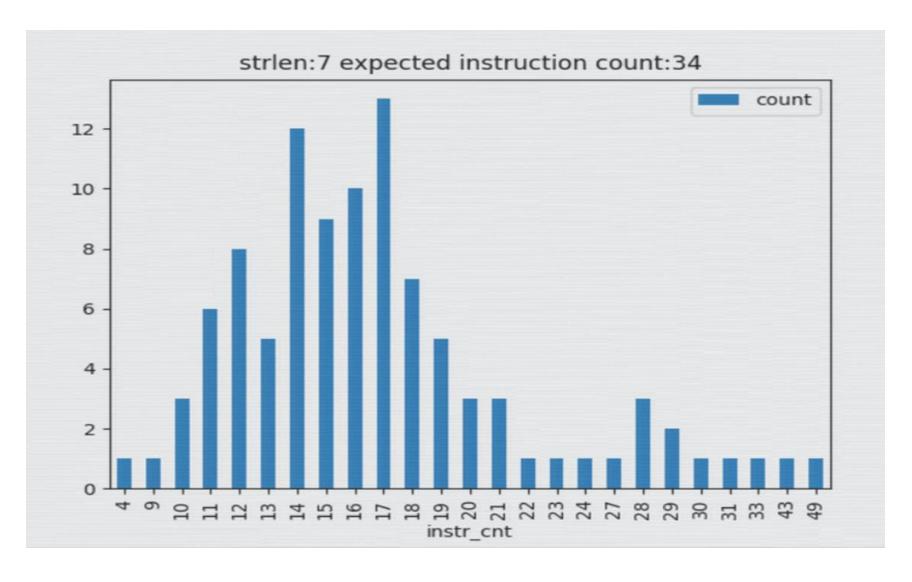


#### Mitigation - Proposal 1



- Mitigation: "Good VMM" portion of KVM implements the stepdetector which will inject #VE (Virtualization Exception) to VM
- VM will eventually exit to the Host

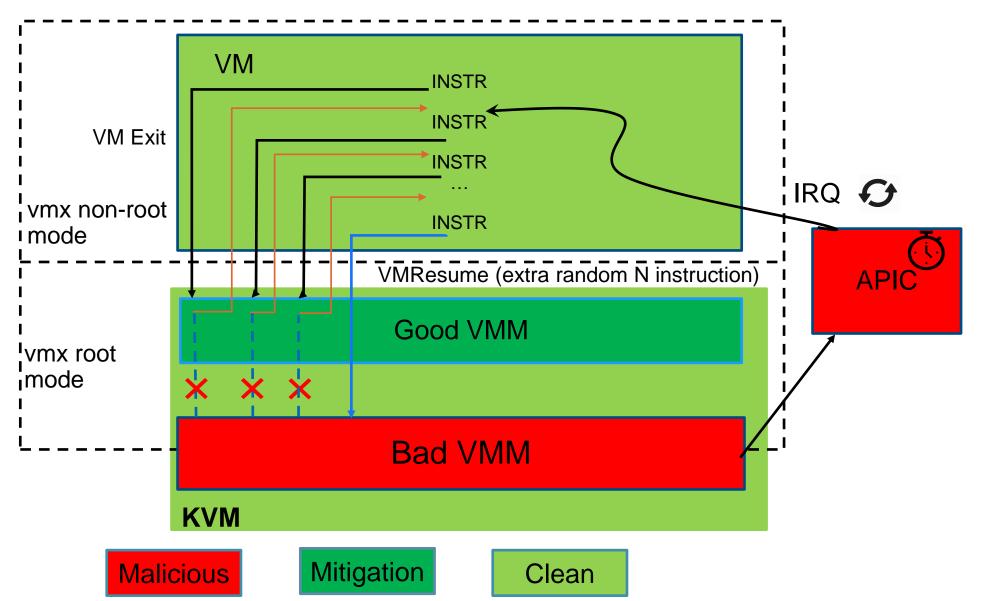
#### **Proposal 1: Mitigation PoC Result**



- Test results show that it wasn't possible to cause VM Exits exactly on every instruction boundary (34 instructions)
- Even if so, the real instruction count of 34 was never met
- Limitations
  - 1) NOT scalable
  - 2) Increased attack surface
- Next step: Define a mitigation to overcome the above limitations



#### Mitigation - Proposal 2



- Mitigation is fully self-contained within the "Good VMM", employing the step-filter
- Resumes the victim VM to execute extra "N" random instructions before it transfer the control back to the Bad VMM

 Result: With the mitigation enabled, instruction count never matched the expected count value of 34, thereby defeating the attack

#### **TDX-Step (v1.0) Mitigation using TDX Module**

```
stepping_filter_e·Vmexit_stepping_filter(
.....vm_vmexit_exit_reason_t·vm_exit_reason,
....vmx_exit_qualification_t·vm_exit_qualification,
....vmx_exit_inter_info_t·vm_exit_inter_info)
{
....tdx_module_local_t*·ld_p·=·get_local_data();
....//·stop·and·reset·EPF·tracking·if·forward·progress·occurred
....uint64_t·guest_rip;
....ia32_vmread(VMX_GUEST_RIP_ENCODE, ·&guest_rip);
....uint64_t·rip_delta·=·vcpu_rip_delta(ld_p, ·guest_rip);
....if·(rip_delta·!=·0)
....{
.....// There·was·forward·progress; ·stop·and·reset·EPF·tracking
.....ld_p->vp_ctx.tdvps->management.last_epf_gpa_list_idx·=·0;
....ld_p->vp_ctx.tdvps->management.possibly_epf_stepping·=·0;
.....ld_p->vp_ctx.tdvps->management.possibly_epf_stepping·=·0;
```

- TDX Module @ https://github.com/intel/tdx-module
- The TDX module employs a VM exit Step-filter algorithm
- Pre-conditions to enable step filter:
  - 1. Interrupt duration should be less 4K cycles since the last entry –TSC (Time Stamp Counter) ) or
  - 2. The RIP of the VCPU has not made progress
- TDX Module resumes the VCPU in "stepping mode" for random number (2-32) of instructions
- VCPU stepping is done using the VMX Monitor Trap Flag (MTF) mechanism.

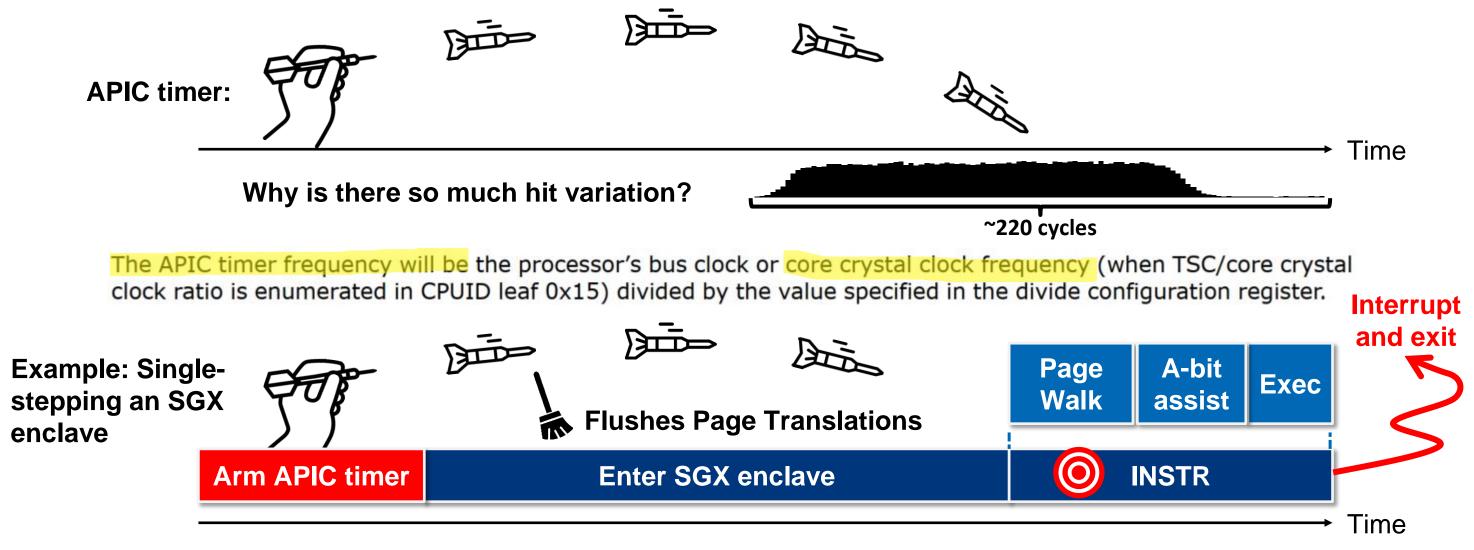


#### **Agenda**

- Intro to Confidential Computing, Intel TDX (Trust Domain Extensions), side-channel Attacks, and malicious single-stepping
- Pre-TDX PoC (Proof of Concept) TDX-step exploit and mitigation
- Techniques to bypass the TDX-Step mitigation, and intro to the new ICSSD (Instruction Counting Single-Step Defense) feature
- Comparison with the SGX-Step mitigation

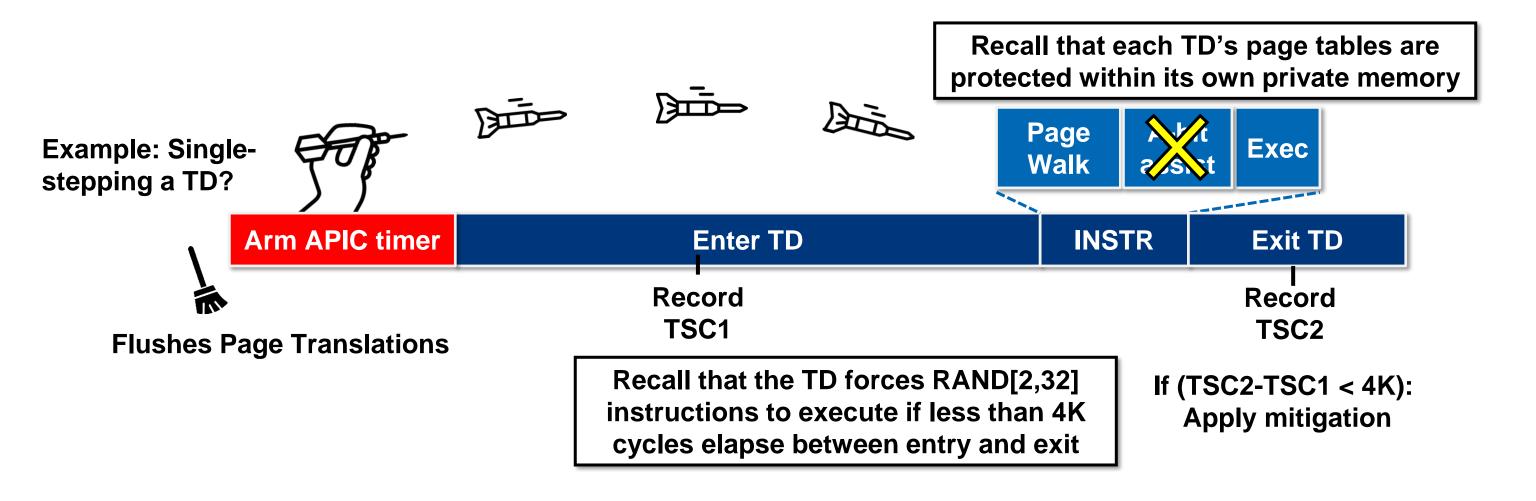


#### (We showed this slide earlier)





#### Attempting to Single-step a TD





#### Demo 1

- Let's attempt to single-step a TD
- (Note that RAX=10 at the start)

#### Listing 1: Evaluation target for single-stepping.

```
62 instructions
  mov gword ptr[r8], 42
loop_label:
                            ~16 instructions
  dec rax
                             forced by the
  nop
                             TDX module
  nop
  nop
  nop
                            ~4 instructions
  jnz loop_label
                               per attack
  mov qword ptr [r9], 42
                                attempt
```

Listing Source: L. Wilke, F. Sieck, and T. Eisenbarth, 'TDXdown: Single-Stepping and Instruction Counting Attacks against Intel TDX', in *Proceedings of the 2024 ACM SIGSAC Conference on Computer and Communications Security, CCS 2024, Salt Lake City, UT, USA, October 14--18, 2024*, 2024.

```
intel@intel-ArcherCity: ~
Dummy unblock for 0x1209c9000 succeeded
Dummy unblock for 0x1209c8000 succeeded
Dummy unblock for 0x1209c6000 succeeded
Dummy unblock for 0x1209c7000 succeeded
Dummy unblock for 0x1209c8000 succeeded
[2024-12-08T05:13:14Z DEBUG apic_attack] Invoking attack start in kernel
[2024-12-08T05:13:14Z DEBUG apic_attack] Triggering victim via network
[2024-12-08T05:13:14Z DEBUG apic_attack] waiting for kernel part to finish
[2024-12-08T05:13:14Z DEBUG apic_attack] checking status
[2024-12-08T05:13:14Z DEBUG apic_attack] Victim is at batch idx 0, state 1, want itera
tions 5. Total idx 15
[2024-12-08T05:13:16Z DEBUG apic_attack] checking status
[2024-12-08T05:13:16Z DEBUG apic_attack] Stopping attack & collecting data
[2024-12-08T05:13:16Z DEBUG tdx_step_userland::freq_sn@ak] outer_len 5
[2024-12-08T05:13:16Z DEBUG tdx_step_userland::freq_sneak] inner_len 7
[2024-12-08T05:13:16Z DEBUG tdx_step_userland::freg_sneak] inner_len 4
[2024-12-08T05:13:16Z DEBUG tdx_step_userland::freq_sn*ak] inner_len 3
[2024-12-08T05:13:16Z DEBUG tdx_step_userland::freq_sneak] inner_len 6
[2024-12-08T05:13:16Z DEBUG tdx_step_userland::freq_sneak] inner_len 5
[2024-12-08T05:13:16Z DEBUG tdx_step_userland::freq_sneak data: outer len
[2024-12-08T05:13:16Z DEBUG apic_attack] done!
intel@intel-ArcherCity:~$
gpa main fn
                : 0x1208f9000
config as cli flags: --ts-target-idx 1 --batch-size 5 --trigger-sequence 0x1209c6000,
0x1209c7000,0x1209c8000,0x1209c9000,0x1209c8000 --allowed-during-attack 0x120907000 -
-stop-gpa 0x1209c8000 --target-code-gpa 0x120907000
Waiting for network package to start...
Waiting for network package to start...
Waiting for network package to start ...
Waiting for network package to start...
Waiting for network package to start...
  0:bash* 1:bash-
                                                     'intel-ArcherCity'
```



#### Single-stepping a TD (TDXdown)

Example: Singlestepping a TD using frequency scaling







Page Walk Exec

Slow victim core frequency

**Arm APIC timer** 

**Enter TD** 

**INSTR** 

**Exit TD** 

TSC1

Intel® 64 and IA-32 Architectures Software Developer's Manual **Example:** Slowing the victim core from 1.8 GHz to 800 MHz will increase the latency of all victim operations by 2.25x, as measured by the TSC

TSC2
TSC2-TSC1 < 4K?

#### 19.17.1 Invariant TSC

The time stamp counter in newer processors may support an enhancement, referred to as invariant TSC. Processor's support for invariant TSC is indicated by CPUID.80000007H:EDX[8].

The invariant TSC will run at a constant rate in all ACPI P-, C-. and T-states. This is the architectural behavior moving forward. On processors with invariant TSC support, the OS may use the TSC for wall clock timer services (instead of ACPI or HPET timers). TSC reads are much more efficient and do not incur the overhead associated with a ring transition or access to a platform resource.



#### Demo 2

 Let's attempt to single-step a TD using frequency scaling (TDXdown technique)

#### Listing 1: Evaluation target for single-stepping.

```
mov qword ptr[r8], 42
loop_label:
    dec rax
    nop
    nop
    nop
    nop
    jnz loop_label
    mov qword ptr [r9], 42
```

Listing Source: L. Wilke, F. Sieck, and T. Eisenbarth, 'TDXdown: Single-Stepping and Instruction Counting Attacks against Intel TDX', in *Proceedings of the 2024 ACM SIGSAC Conference on Computer and Communications Security, CCS 2024, Salt Lake City, UT, USA, October 14--18, 2024, 2024.* 

```
intel® intel-ArcherCity: ~
Dummy unblock for 0x1209c7000 succeeded
Dummy unblock for 0x1209c9000 succeeded
Dummy unblock for 0x1209c8000 succeeded
Dummy unblock for 0x1209c6000 succeeded
Dummy unblock for 0x1209c8000 succeeded
[2024-12-08T05:22:15Z DEBUG apic_attack] Invoking attack start in kernel
[2024-12-08T05:22:15Z DEBUG apic_attack] Triggering victim via network
[2024-12-08T05:22:15Z DEBUG apic_attack] waiting for kernel part to finish
[2024-12-08T05:22:15Z DEBUG apic_attack] checking status
[2024-12-08T05:22:15Z DEBUG apic_attack] Victim is at batch idx 0, state 1, want itera
tions 5. Total idx 15
[2024-12-08T05:22:17Z DEBUG apic_attack] checking status
[2024-12-08T05:22:17Z DEBUG apic_attack] Stopping attack & collecting data
[2024-12-08T05:22:17Z DEBUG tdx_step_userland::freg_sneak] outer_len 5
[2024-12-08T05:22:17Z DEBUG tdx_step_userland::freq_sndak] inner_len 62
[2024-12-08T05:22:17Z DEBUG tdx_step_userland::freq_sndak] inner_len 62
[2024-12-08T05:22:17Z DEBUG tdx_step_userland::freq_sneak] inner_len 62
[2024-12-08T05:22:17Z DEBUG tdx_step_userland::freq_sneak] inner_len 62
[2024-12-08T05:22:17Z DEBUG tdx_step_userland::freq_sneal] inner_len 63
[2024-12-08T05:22:17Z DEBUG tdx_step_userland::freg_sneak] data: outer le
[2024-12-08T05:22:17Z DEBUG apic_attack] done!
intel@intel-ArcherCity:~$
  Avg[]]I
                             0.0% 1792MHz] Tasks: 128, 641 thr, 2173 kthr; 1 running
  Mem[||||
                              7.15G/1007G] Load average: 0.06 0.04 0.07
                                 OK/8.00G] Uptime: 00:55:44
  Swp
                                                                              800MHz]
                                             1[
                                              2[
                                                                        0.6% 1800MHz]
                                            129
                                                                              800MHz]
                                            130
                                                                        0.0% 1800MHz]
                                        F6SortByF7
01 0:bash- 1:bash*
```

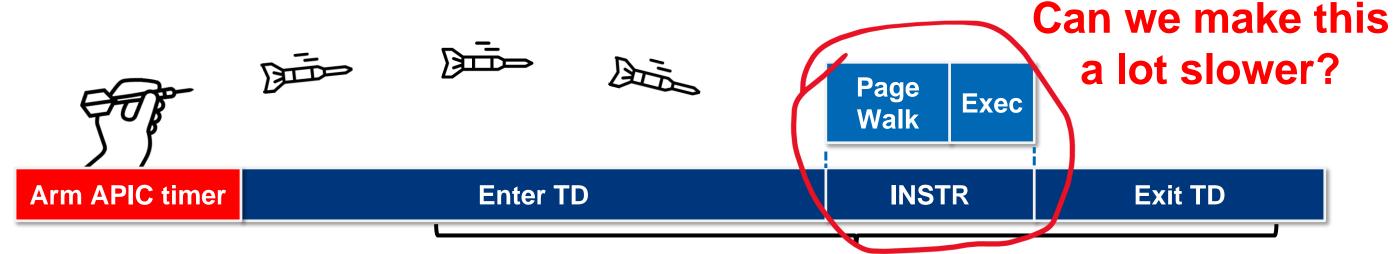


## Key Idea 2: Frequency scaling can be used to fool mitigation heuristics that rely on the TSC

#BHEU @BlackHatEvents



#### Single-stepping a TD (Method 2)



Slowing these operations is the focus of TDXdown

Key Idea 1: Don't try to hit a small target! Instead, make the target *slower* and therefore

## bigger



#### A Page Walk isn't always a Cake Walk

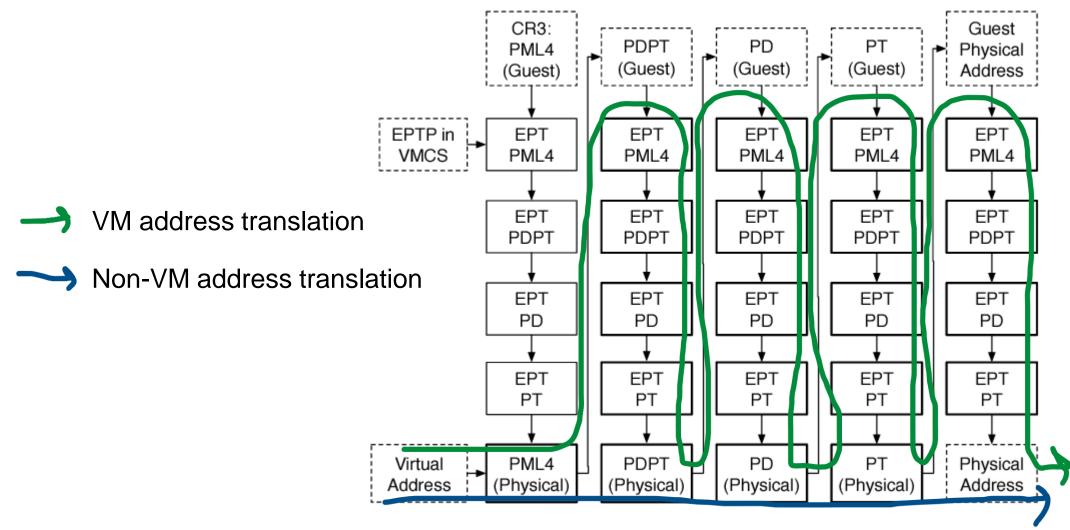


Image Source: https://rayanfam.com/topics/hypervisor-from-scratch-part-4/

#### Demo 3

 Let's attempt to single-step a TD by flushing the EPTs

```
// kmod-wbinvd.c
void apic_callback(void) {
    asm volatile ("\twbinvd\n" ::: "memory");
    fsleep(100000); // 100 ms
```

#### Listing 1: Evaluation target for single-stepping.

```
mov qword ptr[r8], 42
loop_label:
  dec rax
  nop
                 62 instructions
  nop
  nop
  nop
  inz loop_label
  mov qword ptr [r9], 42
```

```
[2024-12-07T06:20:33Z DEBUG apic_attack] Stopping attack & collecting data
[2024-12-07T06:20:33Z DEBUG tdx_step_userland::freq_sneak] outer_len 5
2024-12-07T06:20:33Z DEBUG tdx_step_userland::freq_sneak] inner_len 63
[2024-12-07T06:20:33Z DEBUG tdx_step_userland::freq_sneak] inner_len 64
[2024-12-07T06:20:33Z DEBUG tdx_step_userland::freq_sneak]\inner_len 57
[2024-12-07T06:20:33Z DEBUG tdx_step_userland::freq_sneak] inner_len 65
[2024-12-07T06:20:33Z DEBUG tdx_step_userland::freq_sneak] inner_len 64
[2024-12-07T06:20:33Z DEBUG tdx_step_userland::freq_sneak] data: outer len 5
[2024-12-07T06:20:33Z DEBUG apic_attack] done!
```

Listing Source: L. Wilke, F. Sieck, and T. Eisenbarth, 'TDXdown: Single-Stepping and Instruction Counting Attacks against Intel TDX', in Proceedings of the 2024 ACM SIGSAC Conference on Computer and Communications Security, CCS 2024, Salt Lake City, UT, USA, October *14--18, 2024, 2024.* 



#### Instruction-Count Single-Step Defense

# Intel® Trust Domain Extensions (Intel® TDX) Module Base Architecture Specification

17. Side Channel Attack Mitigation Mechanisms

#### **Suspected Attack Detection Using Instruction Counting**

This attack detection method is applicable if the TDX module implements Instruction-Count Single-Step Defense (ICSSD), as indicated by TDX\_FEATURESO.ICSSD, readable by the host VMM using TDH.SYS.RD\*. It is used only if the TD is not Perfmon-enabled, i.e., ATTRIBUTES.PERFMON is 0. An interruption is considered far enough from the last TD entry if either of the following conditions is true:

- More than one instruction has been retired since the last TD entry, or
- More than one round of a REP-prefixed instruction has been executed since the last TD entry.



#### Demo 4

Let's attempt to single-step a TD when ICSSD is enabled

Listing 1: Evaluation target for single-stepping.

```
62 instructions
  mov gword ptr[r8], 42
loop_label:
                            ~16 instructions
  dec rax
                             forced by the
  nop
                              TDX module
  nop
  nop
  nop
                            ~4 instructions
  jnz loop_label
                               per attack
      qword ptr [r9], 42
                                attempt
```

Listing Source: L. Wilke, F. Sieck, and T. Eisenbarth, 'TDXdown: Single-Stepping and Instruction Counting Attacks against Intel TDX', in *Proceedings of the 2024 ACM SIGSAC Conference on Computer and Communications Security, CCS 2024, Salt Lake City, UT, USA, October 14--18, 2024, 2024.* 

```
intel@intel-ArcherCity: ~
Dummy unblock for 0x11ced1000 succeeded
Dummy unblock for 0x11ced7000 succeeded
Dummy unblock for 0x11cece000 succeeded
Dummy unblock for 0x11ceda000 succeeded
Dummy unblock for 0x11ced7000 succeeded
[2024-12-08T06:01:20Z DEBUG apic_attack] Invoking attack start in kernel
[2024-12-08T06:01:20Z DEBUG apic_attack] Triggering victim via network
[2024-12-08T06:01:20Z DEBUG apic_attack] waiting for kernel part to finish
[2024-12-08T06:01:20Z DEBUG apic_attack] checking status
[2024-12-08T06:01:20Z DEBUG apic_attack] Victim is at batch idx 0, state 1, want itera
tions 5. Total idx 15
[2024-12-08T06:01:22Z DEBUG apic_attack] checking status
[2024-12-08T06:01:22Z DEBUG apic_attack] Stopping attack collecting data
[2024-12-08T06:01:22Z DEBUG tdx_step_userland::freg_sneak] outer_len 5
[2024-12-08T06:01:22Z DEBUG tdx_step_userland::freq_sn(ak) inner_len 4
[2024-12-08T06:01:22Z DEBUG tdx_step_userland::freq_sneak] inner_len 4
[2024-12-08T06:01:22Z DEBUG tdx_step_userland::freq_sneak] inner_len 5
[2024-12-08T06:01:22Z DEBUG tdx_step_userland::freq_snek] inner_len 4
[2024-12-08T06:01:22Z DEBUG tdx_step_userland::freq_sneal] inner_len 6
[2024-12-08T06:01:22Z DEBUG tdx_step_userland::freg_sneak data: outer
[2024-12-08T06:01:22Z DEBUG apic_attack] done!
intel@intel-ArcherCity:~$
  Avg[]
                             0.0% 1792MHz] Tasks: 124, 643 thr, 2183 kthr; 1 running
  Mem[||
                              7.14G/1007G] Load average: 0.02 0.09 0.04
                                 OK/8.00G] Uptime: 01.34.54
  Swp[
                                                                              800MHz]
                                                                        0.0% 1800MHz]
                                            129[
                                                                              800MHz]
                                            130
                                                                        0.0% 1800MHz
[0] 0:sudo- 1:bash*
```

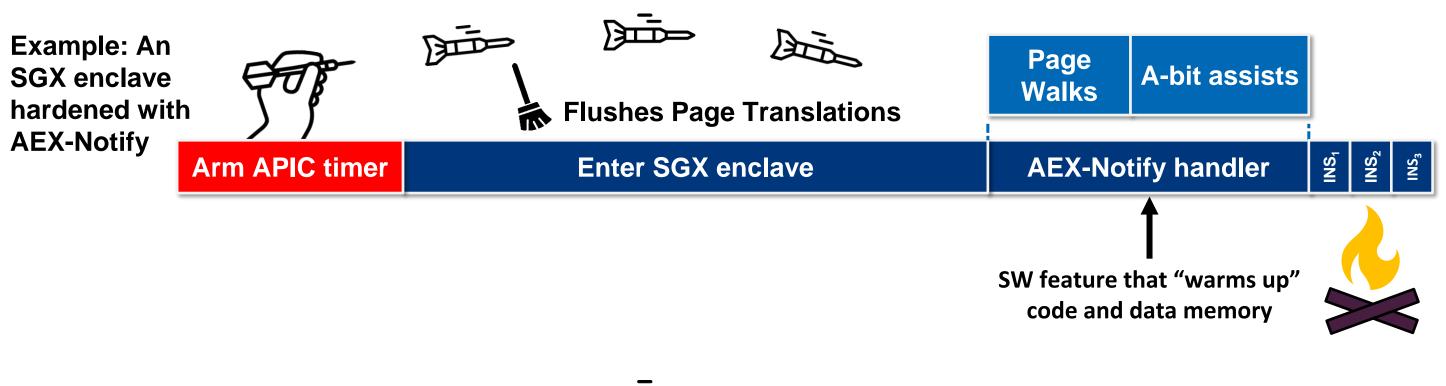


#### **Agenda**

- Intro to Confidential Computing, Intel TDX (Trust Domain Extensions), side-channel Attacks, and malicious single-stepping
- Pre-TDX PoC (Proof of Concept) TDX-step exploit and mitigation
- Techniques to bypass the TDX-Step mitigation, and intro to the new ICSSD (Instruction Counting Single-Step Defense) feature
- Comparison with the SGX-Step mitigation



#### **AEX-Notify: the SGX-Step mitigation**







# Why implement two different single-step mitigations for SGX and TDX?

- TDX is a virtualization-based technology, and therefore has different capabilities
- AEX-Notify uses a SW component that is available in the Intel SGX SDK—a similar SW solution for TDX would require para-virtualization in the guest OS
- Popular OS kernels such as Linux cannot handle arbitrary externally generated events

Trust Boundary: Elements with potential to access confidential data VM Isolation Cloud Stack **BIOS** and Host OS and **VM Guest** Confidential **Guest OS Applications** with Intel® TDX Hypervisor and Admins **Firmware** Admin Data App Isolation Cloud Stack BIOS and Host OS and VM Guest Confidential **Guest OS** Apps and Admins **Firmware** Admin with Intel® SGX Hypervisor Data



#### Conclusion (and Q&A before lunch)

- Confidential Computing's strong adversary model continues to drive exciting research.
   This presentation showed some capabilities that a privileged adversary may use:
  - Arming the APIC timer to hit an adversary-desired instruction with an interrupt
  - Scaling the clock frequency of the victim's CPU core
  - Flushing all of the CPU's caches
- Defense-in-depth mechanisms such as ICSSD and AEX-Notify can help to mitigate potential side-channel attacks against Trusted Execution Environments
- Modern processors use a variety of address translation caching techniques to dramatically accelerate memory accesses in virtualized environments

